

The Context Kit

How to Continue Any AI Project Without Starting Over

By Bogs Copyright © 2026

Table of Contents

Part 1: The Problem

- Chapter 1: When AI Hits the Wall
- Chapter 2: The Context Kit
- Chapter 3: Why This Works

Part 2: The System

- Chapter 4: The Four Files
- Chapter 5: Starting a Session

Part 3: Using It

- Chapter 6: Working Through a Project
- Chapter 7: Managing the Files

Part 4: Keeping It Going

- Chapter 8: Maintenance

Appendices

- Appendix A: Template Library
- Appendix B: Quick Start Guide
- Appendix C: Troubleshooting Guide

Part 1: The Problem

Chapter 1: When AI Hits the Wall

The Problem

You're deep in a conversation with an AI. Maybe you're outlining a novel, planning a business launch, or designing a home renovation. The back-and-forth is flowing. You're making progress.

Then you get the message: *"This conversation is getting long. Some older messages may be dropped."*

Or the chat just stops accepting new messages. It's full.

What happens next? You open a new chat and try to recreate everything. You summarize. You paste notes. You re-explain your project, your decisions, where you were heading.

The AI responds cheerfully: *"I'd be happy to help!"* But it doesn't remember the three ideas you rejected. It suggests things you already tried. It asks questions you already answered.

You're not moving forward anymore. You're repeating yourself.

And here's the real cost: you lose momentum. That nuanced understanding you built up over dozens of messages? Gone. Every time you hit the wall, you're forced to compress hours of thinking into a few paragraphs—and something always gets lost. You end up either oversimplifying (and spending the next twenty messages correcting misunderstandings) or over-explaining (and overwhelming the AI with details it can't prioritize).

Either way, you're not doing the work you sat down to do. You're recreating the work you already did.

Real Example: A Conversation Gone Sideways

Sarah is planning a sustainable coffee shop. Over three sessions, she worked out her target demographic (remote workers, not students), rejected a membership model for year one, chose local roasters over direct trade, and ruled out locations near campus. Session three hits the limit while she's discussing menu pricing.

She opens a new chat: *"I'm planning a sustainable coffee shop for remote workers. Help me with menu pricing."* The AI suggests a student discount program. Then a membership

model. Then direct trade sourcing. All ideas she already rejected. Twenty messages later, she's finally back to pricing—but the AI doesn't remember her \$6.50 ceiling, her three competitor cafes, or her margin assumptions. She's not stuck because the AI is bad. She's stuck because it can't see the path she already walked.

This is the wall. Not the limit itself—but the cycle that follows. The Context Kit process exists to eliminate it. Instead of recreating context, you hand the AI a map.

Chapter 2: The Context Kit

You've hit the wall. Now what?

Most people try one of three things. They write a massive summary and paste it into the new chat—then spend the next ten messages clarifying what they meant. Or they save their conversation history to a document and reference it manually—which means constantly switching between the chat and their notes, losing flow. Or they just start over and hope the AI figures it out from context clues—which it won't.

None of these work. Summaries lose nuance. Documents create friction. Starting fresh wastes time.

The Context Kit is different. Instead of fighting the AI's limitations, you work with them. You organize your project into four simple files, each with one job. When you start a new chat, you upload them in two messages. The AI picks up exactly where you left off—not because it remembers, but because the files tell it everything it needs to know.

The Solution in 30 Seconds

Instead of cramming everything into one conversation, you split your project into four simple text files. Each file has one job. When you start a new chat, you upload them in two messages. The AI picks up exactly where you left off—not because it remembers, but because the files tell it everything it needs to know.

This works with any AI that accepts file uploads—Claude, ChatGPT, Gemini, and others.

That's it. Four files. Two messages. No lost context.

Two Ways to Start

Starting Fresh

You're beginning a new project. Grab the blank file templates from the appendix and upload them to a new chat. Tell the AI what you're working on: *"I'm planning a sustainable coffee shop"* or *"I'm writing a sci-fi novel about AI consciousness."* Then ask it to help populate the files: *"Based on this, help me fill out my Exploration file with the key questions I should be thinking about"* or *"Draft my Execution file with initial guidelines and goals."*

The AI knows the structure. It'll suggest open topics for Exploration, draft guidelines for Execution, and help you establish your Current State in Process. You review, adjust, and save the files. Now you're ready to work.

Mid-Project

You've already got conversations, notes, and decisions scattered across multiple chats. Don't start over. In your current chat—the one with all that context—upload the blank templates from the appendix. Then say: *"I want to organize everything we've discussed into the Context Kit. Help me populate these files based on our conversation history."*

The AI will pull out your decisions for Exploration, suggest guidelines and structure for Execution, and help you capture your current state in Process. You end up with four organized files that preserve all that accumulated context. Save them. Your next session starts with the two-message pattern, and you pick up exactly where you left off.

Either way, the AI does the heavy lifting. You're just guiding what goes where.

What the System Does

Process - Your Current State sits at the top (where you are right now, what's next). Below that: the workflow stages and rules for how you and the AI work together.

Exploration - Open questions you're still thinking through, plus all the decisions you've closed. When you mark something decided, it stays visible with the date and reasoning. The AI stops re-suggesting things you already rejected.

Execution - Three sections in one file: your guidelines (how things should be done), your plan (what you're building), and your action list (what's left to do). Everything the AI needs to stay consistent and on-task.

Progress - A simple log. What got done, what got decided, what's in progress. You add entries as you go. When you come back days or weeks later, you see the path you walked.

You don't upload all four every time. You pick based on what you're doing. The first message always includes Process and Execution—that tells the AI how to work and what's active. The second message brings in Exploration and Progress—that's your context and history.

Why This Beats Monster Prompts

You might be thinking: *"Can't I just write one really detailed prompt?"*

You could. Here's what happens: You write 2,000 words covering your project history, your decisions, your preferences, and your current task. You paste it into a new chat. The AI reads it, responds enthusiastically, and then asks you to clarify something from paragraph four. Or suggests an idea you rejected in paragraph seven. Or focuses on the wrong part entirely because it couldn't tell what mattered right now versus what was just background.

Long prompts create noise, not context. The AI has to guess what's important. Sometimes it guesses wrong.

The Context Kit is different. You're not asking the AI to sort signal from noise—you've already done that by choosing which files to upload. Process says how to work. Execution says what to focus on. Exploration and Progress provide the specific history that matters.

And when the chat fills up? You don't rebuild anything. You open a new chat, paste your Boot Prompt, upload the same four files in two messages, and you're back to work. The AI knows where you are, what you've decided, and what comes next.

It's not about writing more. It's about organizing information so the AI can actually use it.

Before You Hit the Wall

One last thing: before your chat fills up, update your files. Check off completed tasks in Execution. Add what got done to Progress. Update Current State in Process with where you are now. It takes two minutes. Means your next session starts clean instead of trying to remember what just happened.

Chapter 3: Why This Works

How AI Processes Conversations

Here's what happens when you type a message to an AI: it reads your entire conversation history in order, then generates a response. Every word you've written, every response it gives—it processes all of it before answering your new message.

This sounds efficient until you realize the constraint: conversations have a limit. Not because the AI runs out of memory, but because processing longer conversations takes more time and more computing power. At some point, adding another message costs more than it saves. So platforms set a limit—usually measured in thousands of words—and when you hit it, older messages start getting dropped or the chat stops accepting new input.

But here's the important part: *the AI doesn't forget because you run out of space*. It forgets because you stop giving it the information. There's a difference.

When Sarah opens a new chat about her coffee shop, the AI has no access to the previous three sessions. It genuinely doesn't know about the membership model she rejected or the location constraints she's working with. It's not being forgetful—it's starting from zero because that's the only information available.

Why Modular Files Help

The Context Kit works because it separates *how the AI should work* from *what the AI should remember*.

Process and Execution tell the AI *how to work*:

- How you two should interact
- What's in focus right now
- What standards you're following
- What the actual task is

Exploration and Progress tell the AI *what it should remember*:

- What you've already decided (and why those doors are closed)
- What's still open
- What you've already built
- Where the project stands

By splitting them, you're not trying to cram everything into one long message. You're giving the AI exactly what it needs *at the right time*—how to think first, then the specifics to think about.

Compare this to a monster prompt. You write 2,000 words trying to cover: your project goals, your constraints, your rejected ideas, the guidelines you want followed, your current task, and your history so far. The AI reads it all at once and has to guess what matters most right now. Is this message mainly background? Or is the current task buried in there? What can be forgotten after this chat ends, and what needs to be preserved?

The AI doesn't know. So it plays it safe and remembers everything—which means it repeats things you already covered, suggests ideas you already rejected, and asks questions you already answered.

With modular files, the AI doesn't have to guess. Process says *"here's what we're doing right now."* Execution says *"here's what you need to focus on."* The AI isn't sorting through a wall of text trying to figure out priorities—the priorities are already sorted.

The Cost of Cramming Everything In

This might seem like a small difference, but it compounds.

A monster prompt costs you in two ways:

First, it makes the AI less precise. When you dump 2,000 words into a message, the AI has to guess what's important. It sees your rejected ideas and treats them as context—which is right—but then when you ask a question, it might suggest one of those ideas anyway because it doesn't have clear boundaries between "background information" and "active constraints." You end up correcting it, which means more messages, which means you hit the wall faster.

Second, it makes every chat harder to start. When you open a new conversation, you need to paste the whole thing again. Same 2,000 words. If you have refined your understanding since you last pasted it, you have to update the whole document. If you only want to reference certain parts, you still paste the whole thing because extracting just the relevant sections takes time.

The Context Kit avoids both costs. By splitting files by function, the AI knows exactly what to prioritize. And when you start a new chat, you only upload what you need right now—which is faster, clearer, and less prone to the AI misinterpreting your constraints.

What This Means For Your Work

The system works because it respects how conversations actually function. AI doesn't magically synthesize complexity—it processes information sequentially and uses whatever's available to generate a response. The Context Kit doesn't fight that. It works *with* it.

You're not outsmarting the AI or optimizing some hidden setting. You're just being explicit about what matters, when it matters, so the AI can focus instead of searching.

This is why you can come back days or weeks later, upload the same four files, and pick up exactly where you left off. You're not recreating context. You're handing the AI a map that says: here's where we are, here's what we've decided, here's what we're doing next.

The AI can work from that immediately. Because the files say what matters.

Part 2: The System

Chapter 4: The Four Files

Here's what the system looks like in action:

You're planning a coffee shop. Your **Process** file says you're in Execution, working on menu pricing. Your **Execution** file has your pricing guidelines (35% margin minimum, \$4.50-6.50 competitive range) and shows the pricing model as the active task. Your **Exploration** file shows you already decided against student discounts and membership models—those doors are closed. Your **Progress** file shows you spent two sessions on location research and one session finalizing your target demographic. You upload all four files to a new chat. The AI immediately knows: where you are (menu pricing), what standards to follow (margin and price range), what not to suggest (discounts, memberships), and what you've already completed (location, demographics). You start working on pricing tiers without re-explaining anything.

That's the system. Four files, each with one job. Let's break down what each file does and why it matters.

The Process File: How You Work Together

This file is your north star. It tells the AI where you are, what you're focused on, and how to work with you.

What goes in it:

- Workflow stages (how you move through the work)
- Rules for the AI (what to prioritize, what not to change)
- Current State (where you are right now)

Why it matters: When you upload this file, the AI immediately knows where you are and what's active. No guessing. You're telling it directly: *"We're in Execution, working on Chapter 4. Everything before this is done."*

What it looks like:

```
**Phase:** Execution
**Last Completed:** Chapter 3 drafted
**Next:** Draft Chapter 4
```

That's enough. The AI knows exactly what to focus on.

The Exploration File: What You've Decided

This file tracks open questions and closed decisions. It's your decision log.

What goes in it:

- Open Topics: questions you're still exploring, options you're weighing
- Closed Decisions: what you decided, why, and when

Why it matters: This is how the AI learns your closed doors. Sarah's coffee shop decided against a membership model for year one. When she uploads this file, the AI stops suggesting memberships. It knows that the door is closed.

Decisions stay visible with dates and reasoning. Weeks later, if you wonder whether to revisit something, you can see why you ruled it out originally.

What it looks like:

```
**[2025-01-09] Target Audience**  
- Remote workers, not students  
- Better margins, consistent presence  
- Can add student pricing later if needed
```

The Execution File: How You'll Do It

This file is your playbook. It contains guidelines, your plan, and what's left to do.

What goes in it:

- Guidelines: style, tone, constraints, quality standards
- Plan: what you're building, structure, scope, what done looks like
- Action: current focus and task list

Why it matters: This keeps you consistent. If your guideline says "examples under 3 paragraphs," every section follows that. If your plan says "8 chapters," the AI won't suggest chapter 9. The action list shows what's done and what's next.

What it looks like:

```
Guidelines:  
- Clear, conversational  
- Examples under 3 paragraphs
```

- No jargon

Action:

- [✓] Chapter 1: When AI Hits the Wall
- [✓] Chapter 2: The Context Kit
- [] Chapter 3: Why This Works (current)

The AI sees what to work on, what standards to follow, and what's already handled.

The Progress File: What You've Done

This file is an append-only log. Add entries as you work. It's your project history.

What goes in it:

- Date-stamped entries
- What got completed
- Decisions made during the session
- What's next

Why it matters: When you return days or weeks later, you see the entire path you walked. Not summaries—actual milestones. This keeps you oriented and creates a record of momentum. You can see what took one session versus three, when you changed direction, and why.

What it looks like:

```
## 2025-01-09 - Chapter 3 Drafting
**Completed:** Chapter 3 all sections drafted
  - How AI processes conversations
  - Why modular files help
  - Cost of cramming everything in one chat
```

How They Work Together

You don't upload all four files every time. You're strategic.

First message (when resuming): Upload Process + Execution

- Process: how to work, where you are

- Execution: what's in focus, what standards to follow
- Together: your working context

Second message (when you need history): Upload Exploration + Progress

- Exploration: what's decided, what doors are closed
- Progress: the path you walked
- Together: your decision context and project history

Use both messages every session, or just the first if you're moving forward and don't need past decisions. The system flexes.

Creating Your Files

Appendix A has the complete blank templates ready to copy. The next chapter walks you through starting your first session—either fresh or by organizing work you've already started.

For now, just understand what each file does and how they divide the work. Four files, four purposes. Everything in the right place so the AI knows what matters and you can pick up weeks later exactly where you left off.

Chapter 5: Starting a Session

You know what the four files do. This chapter explains how to use them—specifically, how to start a session, what to upload, and in what order.

The answer is the two-message pattern.

The Two-Message Pattern Explained

When you resume work on a project, you send two messages in a new chat.

Message 1: Upload Process + Execution. Tell the AI how you work and what you're focused on.

Message 2: Upload Exploration + Progress. Give the AI your decision context and history.

Why split it? Chat platforms have upload limits. More importantly: this order makes sense. You orient the AI to how you work *before* giving it your history. Process and Execution are active, operational files. Exploration and Progress are context—the backstory.

It takes 30 seconds.

First Message: Setting Up the Session

This is your "System Boot" message. Here's the prompt:

```
I'm working on a project using a file system to keep context between sessions. I'll share files in two messages:

**This message:** Process (how we work + where we are) + Execution (guidelines + plan + tasks)
**Next message:** Exploration (what's been decided) + Progress (what's been done)

Here's what I need from you:
- Focus only on what's in Current State
- Don't reopen closed decisions unless I ask
- Follow the guidelines and quality standards in Execution
- Help me update files at the end so I can pick this up later

I'll paste the first two files now.
```

Then paste Process and Execution.

The AI reads this and understands: *Here's how we work. Here's where we are. Focus on the Current State.* It's a clear handoff that sets boundaries—the AI knows to focus on your active work, not second-guess closed decisions.

Second Message: Loading Your Context

In your next message (same chat), paste Exploration and Progress. Brief note:

Here's the context from earlier work.

That's it. The AI now has the full picture: how to work, where to focus, what's been decided, what you've already done.

Understanding the File Numbers

You might notice something: you upload files 1 and 3 in the first message, then files 2 and 4 in the second. Why skip file 2?

The numbering follows your workflow sequence, not upload order:

- **01-process** = How you work (process comes first)
- **02-exploration** = What you're exploring (exploration happens next)
- **03-execution** = What you're doing (execution follows exploration)
- **04-progress** = What you've done (progress tracks it all)

But the upload grouping is about function:

- **Message 1 (files 1+3):** Active work files—Process and Execution tell the AI how to work and what's in focus right now
- **Message 2 (files 2+4):** History files—Exploration and Progress provide context and decisions from past work

The numbers reflect how you think about the work. The upload pattern reflects what the AI needs when.

It's not random—it's intentional. Workflow sequence for the files, functional grouping for the upload.

Optional: Session Prompts for Consistency

If you work on a project regularly, you might find yourself typing the same boot prompt every session. There's a better way.

Create a file called `00-prompts.md` in your project folder. Save your standard session starters there: the boot prompt you use to open sessions, the end-of-session update request, even a quick context reset prompt if the AI starts drifting. When you need to start a session, copy from your prompts file instead of retyping.

This isn't required—the system works fine without it. But if you're working on something for weeks or months, having your prompts ready to copy saves time and keeps your sessions consistent. See Appendix A for a template you can customize for your project.

Example: Resuming After a Week

Sarah's been working on her coffee shop plan. Last week: menu pricing and staffing. Now she's taking a week off, then returning to scout locations.

She opens a new chat. First message:

I'm working on a project using a file system to keep context between sessions.
I'll share files in two messages.

[... boot prompt ...]

Here are my Process and Execution files.

She pastes those two files.

The AI reads them. Current State says: "Last Completed: Menu & Staffing. Next: Location Scout." The Execution file shows her location criteria and checklist. The AI knows exactly where she left off.

Second message:

Here's my full context.

She pastes Exploration and Progress. The AI sees every decision: demographics, no membership model, local roasters, margin targets. The Progress file shows it took three sessions to nail pricing but only one for staffing.

Now the AI can help her scout locations *in the context of everything she's decided*. It won't suggest direct trade sourcing (already ruled out) or ask about student discounts (already closed).

Different example: Marcus is building a SaaS product. He worked on the pricing model last session. This week he's resuming to design the onboarding flow. Same pattern: upload Process + Execution (shows current focus: onboarding design), then Exploration + Progress

(shows closed decisions on pricing tiers, target users, competitive positioning). The AI immediately knows what's in scope and what doors are closed.

Starting Fresh

Beginning a brand new project? Same pattern, but start with the AI's help.

First message: boot prompt + blank templates (from Appendix A). Then ask:

Based on what I just told you about my project, help me populate these files.

I need you to draft initial Exploration questions, suggest guidelines for Execution, and help me define my Current State and goals.

The AI helps you fill them out. You review, adjust, save. Now you're ready to work with the full system from session one.

Session Patterns for Different Work

Not every session needs both messages.

Deep work (drafting, planning, iterating): Use both messages. You need full context.

Quick check (reviewing a closed decision, adding a task): Use Message 1 only. You don't need a full history.

Brainstorming (exploring new features, considering changes): Use both messages. You'll want to know what you've already decided before opening new territory.

The pattern flexes. The rule: always upload Process and Execution, because they're active. Exploration and Progress are optional based on what you need.

But most of the time? Both messages. Takes 30 seconds and keeps everything aligned.

You've loaded your files and started your session. Chapter 6 covers what happens next: how to work through a project, when to update files, and how to maintain momentum.

Part 3: Using It

Chapter 6: Working Through a Project

You've started your session (Chapter 5 covered the boot pattern). Process and Execution are uploaded. Exploration and Progress are loaded. You and the AI are working.

This chapter covers the rhythm of active work: moving through phases, updating files, handling decisions, and staying focused.

Moving Through Phases Naturally

Your workflow phases live in Process. You might be in Exploration (figuring things out), Execution (building), or Progress (consolidating). As you work, you naturally move through them.

In Exploration, you're asking questions: "Should we offer a student discount?" "What's our profit margin target?" The AI helps you think. You're not deciding yet—you're exploring.

When a question gets answered, move it to Exploration's Closed Decisions section. Date it, explain why, move on. Now it's off the table.

In Execution, you're doing the work. Following guidelines. Checking off tasks. Pushing forward.

You don't have to declare "We're moving to Execution now." It happens naturally. The file structure keeps track of where you are without slowing you down.

When to Update Files

Update at session endpoints, not continuously.

End of session: Before closing the chat, ask the AI to update your files. Takes 1-2 minutes. The AI generates updated markdown for all four files based on what you accomplished. You review, copy, save.

The prompt:

Based on this session, please update my four files:

- Check off completed tasks in Execution
- Move any closed decisions to Exploration with dates
- Add a Progress entry for what we accomplished

– Update Current State in Process to reflect where we are now

The AI generates the updates. You review them (make sure they're accurate), then save the files.

During a session: Only update if something structurally shifts. Closed decision needs reopening? Tell the AI to update Exploration. New constraint changes the plan? Update Execution. Routine task checking? Save it for the end.

The rhythm: Start chat → Upload files → Work 1-3 hours → Ask AI to update files (1-2 min) → Review and save → Done.

When to edit manually: If the AI gets something wrong—missed a task, wrong date, unclear wording—fix it yourself. You understand what each file needs (Chapter 4 covered that). The AI is your productivity tool, not a replacement for judgment. Quick manual edits take seconds.

Chapter 8 covers the detailed update process and what to verify before saving.

When a Closed Decision Needs to Reopen

This happens. You're writing a novel and decided the protagonist doesn't have siblings. Two chapters in, you realize the sibling dynamic is essential for the conflict.

Tell the AI: *"We need to reopen the sibling decision. Move it from Closed Decisions back to Open Topics."*

You explore it again from the new context. When you decide, tell the AI to move it back to Closed Decisions with updated reasoning: *"Reopened because of conflict dynamics. Now: protagonist has one sibling."*

It's tracked. Your Progress log shows it happened. Your next session knows the decision shifted and why.

When to Add or Refine Guidelines

As you work, you might need tighter guidance. Your tone guideline says "conversational" but you're drifting between casual and formal. Or "examples under 2-3 paragraphs" but you're naturally writing longer.

Mid-session, call it out: *"I'm noticing my examples are running longer. Should we adjust the guideline or tighten my writing?"*

Work it out. Then tell the AI to update Execution with the refined guideline. If "conversational" becomes "conversational but not colloquial; friendly expert tone," the AI updates it. The guideline gets better with use.

At session end, the Progress log will note: "Refined guideline for tone."

Staying Focused vs. Distracted

Here's where the Current State in Process matters most.

Your Process file says: "Next: Draft Chapter 5." You and the AI are working on Chapter 5. Then you think of an interesting example for Chapter 3. Or you realize Chapter 4 might need restructuring.

Two choices:

Stay focused: Tell the AI to log it in Execution: "Future consideration: Chapter 4 restructure idea." Keep working on Chapter 5. Deal with the idea later.

Go distracted: Pivot to Chapter 4. Spend 30 minutes on the idea. Your session becomes unfocused. You end with Chapter 5 unfinished and Chapter 4 half-explored. Your next session is confused about what's actually in progress.

Different context: You're building a product feature. Current State says "Next: Design onboarding flow." Mid-session, you realize the pricing page needs a quick fix. Log it: "Future: Pricing page adjustment needed." Finish onboarding. Fix pricing later when it's the active focus.

The system doesn't prevent distraction—but it makes it visible. If you pivot, update Current State. Your next session sees: "Last Completed: Partial Chapter 5 + Chapter 4 exploration. Next: Decide whether to continue Chapter 5 or pursue Chapter 4 restructure."

It's a choice, not a mess.

Dealing with Feedback from Outside

Sometimes you get feedback during a session. A friend reviews your work. Your boss has notes. A client suggests a change.

When this happens:

- Don't immediately incorporate it
- Tell the AI to add it to Exploration as an open topic: "Feedback: Friend suggested adding X to Chapter 2. Does this align with goals?"
- Discuss it with the AI if significant
- Update Closed Decisions when you've decided
- Move on

Example from consulting work: A client suggests adding a new deliverable mid-project. Instead of saying yes immediately, have the AI log it in Exploration: "*Client request: Add competitive analysis to Phase 2. Consider: Does this fit the timeline? Does it align with scope?*" Evaluate it against your closed decisions (project scope, timeline constraints) before committing.

This prevents reactive pivots. External feedback stays in a structured place where you can evaluate it against your decisions.

The Goal: Momentum

The point of this system isn't rigidity. It's momentum.

You want to come back to a project and immediately know where you are, what's done, and what comes next. No ramp-up time. No "*wait, did I finish Chapter 5?*" checking.

The files do that work. They're the external record of your thinking, so your brain stays free to think.

Chapter 7: Managing the Files

You know what the files are. You know how to start a session. You know how to work. Now: practical questions. What if a file gets really long? What if you want to focus on just one task? When do you start a completely fresh chat instead of resuming?

This chapter is about file maintenance during active projects.

File Rotation Strategy

Over time, files grow. Process stays small, but Execution gets longer (more tasks, more guidelines). Progress becomes a long log. Exploration accumulates decisions.

This is normal. But there's a tipping point: when files get so long that uploading them slows things down or makes it hard to find what matters right now.

When this happens, you rotate files.

What that means: You start fresh versions of some files, keeping only what's active. You archive the old versions.

Example: Your Progress file has 30 entries from a month of work. You're now in a new phase. You create a new Progress file with just today's entry, and rename the old one to "progress-archive-december.md". Your archive is safe if you need to reference it, but your active chat isn't bogged down by a long history.

Similarly with Execution: If your task list grows to 50 items, with most checked off, you can archive the old list and create a new one with just active tasks. The archive shows what you've done; the active file shows what's next.

Exploration is trickier—don't archive closed decisions unless you're sure you'll never reference them. But as a project matures, you might move truly old decisions to an archive. Keep recent decisions active.

Process is small enough that you rarely need to rotate it.

The rule: When a file gets unwieldy, split it. Keep active content in the main file, archive history separately. Your current chat only needs the active file.

Swapping Files As You Progress

Some projects naturally have phases. You might plan something for three sessions, then execute it for five sessions.

During planning, you might upload all four files every session. You're exploring, so Exploration and Progress are both valuable contexts.

During execution, you might upload just Process and Execution. You know what you decided; you don't need to re-read all the exploration. Your focus is: here's where we are, here's what to do.

This isn't a rule—it's a flexibility. Some sessions you need all four. Some sessions you only need one or two.

The system adapts to how you work. You're not locked into a pattern. You're using the files that serve the moment.

When to Start a Fresh Chat

Sometimes you don't resume—you start over. This usually happens for one of three reasons:

1. The previous chat filled up (or got too long).

You were working on Chapter 5. The chat got long. Rather than resume the same chat and deal with limits, you start a new chat. Upload your files using the two-message pattern. Keep going with Chapter 5. It's a hard reset, but your context is preserved in the files.

2. You're pivoting significantly.

You've been working on a memoir but decided to pivot to an essay. Your Execution file changes. Your Current State changes. Rather than clutter the previous chat with that shift, you start fresh. Upload the updated files, and the AI knows you're in a new phase.

3. You're context-switching to a different project.

You've been working on Project A. Now you need to work on Project B. Upload Project B's files, and the AI switches context. When you come back to Project A, you upload its files again. Fresh chat, full context, no confusion.

Multi-Project Workflow

If you manage multiple projects, the system scales easily.

For Project A: Keep files in folder `project-a/kit/` with names `01-process.md`, etc.

For Project B: Keep files in folder `project-b/kit/` with names `01-process.md`, etc.

When you switch projects, you switch file sets. Each project has its own four files, isolated.

You can work on Project A on Monday (two-message pattern with A's files), then switch to Project B on Tuesday (new chat, B's files), then back to A on Wednesday (new chat, A's files). The AI never mixes them up because the files are explicit.

Managing File Size Over Months

If you're working on a long-term project (months, not weeks), files will grow. Here's a practical approach:

Progress: Archive quarterly. Keep the last month active, move older entries to an archive file.

Execution: Prune checked tasks quarterly. Move "done" sections to an archive. Keep only active tasks and future phases.

Exploration: Keep all closed decisions indefinitely (they're your decision history), but you can move truly old ones to an archive if the file gets unwieldy.

Process: Almost never needs archiving. It stays small.

This keeps your active files lean while preserving history.

When to Kill a Project File

Eventually, a project ends. You finish the book, launch the product, complete the plan.

Your last session: update Progress with the completion. Save your files. You're done.

You can keep the files indefinitely (they're good reference material), but you stop uploading them to new chats. The project is closed.

If you ever need to revisit it—a sequel, a Phase 2, lessons learned—you have the files with full context. Upload them, update Current State, and you're back in the project.

The Key: Simplicity

Managing files sounds complex, but it's not. The defaults work for most projects:

- Start with the four files
- Update them at session end
- Rotate files only when they get long
- Switch projects by switching file sets

That's it. You don't need elaborate file management. You're not running a database. You're keeping track of context so you can work more effectively.

The files serve you. When they stop serving (because they're too long or out of date), you fix them. Otherwise, you don't think about them. You think about your work.

Part 4: Keeping It Going

Chapter 8: Maintenance

Every project needs maintenance. Not because things break, but because momentum is fragile. This final chapter covers long-term practices: how to end sessions properly, keep files useful, and sustain projects over time.

Ending Sessions Properly

Before you close a chat, have the AI update your files. Takes 1-2 minutes and keeps your next session seamless.

The prompt:

Based on this session, please update my four files to reflect what we accomplished.
Generate the updated markdown for all four files.

The AI will generate updates for:

1. Execution: Checks off completed tasks. Marks anything in-progress. Adds new tasks if they come up.

2. Exploration: Moves closed decisions from Open Topics to Closed Decisions with dates and reasoning. Adds new open questions if they emerge.

3. Progress: Adds an entry for today with:

- What got completed
- What decisions got made
- What's in progress or next
- Anything else worth remembering

4. Process: Updates Current State:

- Last Completed: What was the most recent milestone?
- Next: What comes after that?

Review before saving: Read through the AI's updates. Make sure:

- Tasks are accurately marked (did you really finish that section?)
- Decisions have correct dates and clear reasoning

- Progress entry captures what actually happened
- Current State reflects where you actually are

If something's wrong, tell the AI to fix it. Or edit the file yourself—you know what each file needs (Chapter 4 covered that structure).

Manual updates when needed: Sometimes you need quick manual edits:

- AI missed a small task completion
- Wording in a decision isn't quite right
- Progress entry needs a detail added
- Current State needs minor adjustment

Open the file, make the edit, save. Takes 30 seconds. The AI handles the bulk of maintenance, you handle precision where needed.

Why this matters: Projects die when momentum dies. Momentum dies when you come back after three days and can't remember if you finished the pricing section or not. You update files so you never have that problem.

Keeping Files Lean

Over time, files accumulate content. Tasks get done. Decisions pile up. Progress entries multiply.

Lean files serve you. Bloated files slow you down.

Here's what lean means:

Execution: Only active and upcoming tasks. If you've done 20 tasks and have 5 left, check off the 20 and keep just the 5 visible. Archive the completed section to a separate file if you want the history, but your active file stays focused.

Progress: The last month of entries visible, minimum. If you're three months into a project, you can archive the first two months to a separate file. Your active progress log shows recent momentum, not ancient history.

Exploration: Keep closed decisions (they're your reasoning), but if you have 50 closed decisions and only 5 of them are relevant to your current work, you can move the old ones to a reference archive. The active Exploration file is your current-project context.

Process: This is already small. Almost never needs pruning.

The rule: **If an item doesn't affect today's work, move it to an archive.**

Asking the AI to clean files: When files get bloated, ask the AI:

My Execution file has 30 completed tasks cluttering it.

Please generate a cleaned version with only active tasks,
and a separate archive file with the completed tasks for reference.

The AI creates both files. You save them. Your active file is lean again.

Archives aren't garbage—they're history. You keep them. But your active files are clean and purposeful.

Guarding Against Context Rot

AI isn't a person; it's a statistical engine with a "sliding window" of attention. As your session grows long, the AI's performance can degrade—a phenomenon we call **Context Rot**. You'll notice it when the AI starts ignoring your **Current State**, giving you broken formatting, or "forgetting" decisions you closed days ago.

Think of it like a conversation at a loud party: eventually, the AI loses the thread. When this happens, don't keep pushing. Instead, you need to perform a **Context Reset** to clear the "noise" and refocus the AI on your Kit files.

Note: For the full step-by-step recovery protocol to save your progress and reset the session, see **Appendix C: Troubleshooting Guide**.

Sustaining Momentum Over Months

Long projects are won through consistency, not heroic sprints.

The system is designed for this. You upload the same four files week after week, month after month. The AI immediately understands where you are because Current State is explicit. Closed decisions don't get re-suggested because Exploration is explicit. You don't re-explain the project because Execution is explicit.

This eliminates friction. Every session, you save 10-15 minutes that would've been spent on context recreation. Over a month-long project, that's hours of extra productive time.

But there are practices that keep momentum alive:

1. Consistent session structure

Same pattern every time: upload files (2 min), brief orientation (1 min), work (45 min - 2 hours), AI updates files (1-2 min). It becomes a ritual. Your brain knows what to expect.

2. Visible progress

Your Progress log is an artifact. When you're tired and wondering if you're making progress, you can read it and see: "Three weeks ago, I drafted an outline. Two weeks ago, I drafted the first section. Last week, I finished section 2. This week, 30% through section 3." The progress is real and visible.

3. Regular file review

Once a week, spend 5 minutes reviewing your Process and Execution files. Does Current State still match reality? Are the guidelines still accurate? Does the plan still make sense? Small adjustments keep the system aligned.

4. Knowing what to do next

Your Execution file always has a "Next" section. When you sit down, you don't stare at a blank screen wondering what to work on. You open Execution, and the next task is there. This sounds trivial, but it's crucial. Starting is the hardest part of the work. Making starting obvious is everything.

Dealing with Long Breaks

Sometimes you have to stop. Vacation. Emergency. A competing project. Life happens.

When you come back after a week, or a month, your files become invaluable.

Your Process file says: *"Last Completed: Initial menu draft. Next: Pricing model analysis."*

Your Progress file shows the last entry: *"2025-01-15: Completed menu draft, identified three pricing models to analyze, selected Excel as pricing tool."*

Your Exploration file shows the open question: *"How much should we charge per drink? Analyzing three models: cost-plus (35% margin), competition-based (\$4.50-6.50 range), value-based (survey customer willingness)."*

You open a new chat, upload your files, and you're back in the project immediately. No ramp-up time. No "What was I working on?" No lost momentum. You pick up exactly where you left off, even if it's been three weeks.

When a Project Gets Stuck

Some projects stall. You planned to work on it every day, but it's been two weeks and you haven't touched it.

Stalled projects don't restart themselves. They need a gentle push.

Here's what works: Open a new chat. Upload your files. Don't try to jump back into execution. Instead, spend 30 minutes reviewing:

- Read your Progress log from the last month. Remember what you've done.
- Review Exploration. What are you stuck on? Is there an open question that's blocking progress?
- Check Execution. Is the next task clear, or is it too vague?

Often, the stall isn't lack of motivation—it's unclear next steps or an unresolved decision blocking progress.

Once you clarify what's stuck, ask the AI: *"What's the simplest next step to get unstuck?"* Work on that one small thing. Have the AI add it to Progress. Often, that small progress is enough to restart momentum.

Projects That Actually End

Some projects finish. You complete the book. You launch the business. You finish planning the renovation.

When a project ends, ask the AI to make final updates:

This project is complete. Please update the files to reflect:

- Final Progress entry: "Project complete. All 8 chapters were drafted and polished."
- Current State: "Phase: Complete. Next: Archive files and close project."

Save your files. You can store them indefinitely. Months or years later, if you need to revisit the project, reference it, or do Phase 2, the files are there with complete context.

But you stop uploading them to new chats. The project is closed.

When Files Get Long: Archive Strategy

Over time, your files accumulate history. Progress entries multiply. Execution tasks pile up. Closed decisions grow. This is normal—it means you're making progress.

But there's a tipping point. When files get so long that uploading them slows down sessions or makes it hard to find what matters right now, it's time to archive.

Archiving means splitting files: keep active content in your main files, move old history to separate archive files. Your Progress file might become `04-progress.md` (current month) and `archive/2025-01-progress.md` (January entries). Your Execution file keeps active tasks, archives completed ones to `archive/2025-01-completed-tasks.md`.

The pattern: **active content stays in main files, history moves to dated archives**. You can ask the AI to help create archives—it'll extract old content, generate the archive file with proper dating, and give you a clean active file to review and save.

Archives aren't garbage—they're reference material. When you need to remember why you made a decision three months ago or see how long a phase took, the archives have that history. But your daily work uses lean, focused files.

For detailed archive organization patterns, naming conventions, and file structure examples, see Appendix C: Advanced Archive Management.

Project Files vs Kit Files: A Clarification

There's an intentional relationship between the files in `kit/` and the project files you create.

Kit files are templates, not instances.

The `kit/` directory contains:

- `01-process.md` - a blank template showing what a Process file looks like
- `02-exploration.md` - a blank Exploration template
- `03-execution.md` - a blank Execution template
- `04-progress.md` - a blank Progress template

When you start a project, you:

1. Copy these templates to your project folder
2. Ask the AI to help you fill them in with your specific context
3. Evolve them as your project grows

The kit files never change. They're the reference. Your project files are instances of the kit pattern.

This distinction matters for three reasons:

1. **Consistency:** Every project follows the same four-file structure, making it easy to switch between projects
2. **Reusability:** When you finish one project and start another, you already know the pattern
3. **Reference:** If a project file gets messy, you can check the kit files to remember the intended structure

Practical relationship:

When your project's Execution file becomes cluttered, compare it to `kit/03-execution.md`. Does your structure still match the template? Are the sections still clear? Sometimes, seeing the template reminds you to clean up your version.

This is especially useful for Exploration files. The kit version shows the pattern: "Open Topics, Closed Decisions, with date + reasoning." If your project version has drifted, realigning to the template makes it more useful.

You can extend the template, but don't break it. If you need a new section (say, "Constraints" in Execution), add it—but keep the core structure intact. This preserves the pattern across projects.

Adding Variations to Examples

Examples are powerful, but a single example can feel limiting. "But what if my project is different?" The answer: most examples adapt.

The Coffee Shop Example Pattern:

Throughout this book, Sarah's coffee shop example appears. Let's show how the same framework works for different contexts—and how you'd adapt it.

Same four files, different contexts:

1. Coffee Shop Startup (Sarah)

- Process: Exploration of business model, then execution of menu/pricing
- Execution tasks: Define pricing model, design menu, identify target audience
- Exploration: Market positioning, customer demographics, profitability target

2. Freelance Consulting Business

- Process: Exploration of service offerings, then execution of website/rate structure
- Execution tasks: Define service packages, write case studies, set pricing
- Exploration: Target industries, rate benchmarks, positioning vs competitors

3. Product Launch (Software)

- Process: Exploration of market fit, then execution of MVP features
- Execution tasks: Build core features, design UI, plan launch sequence
- Exploration: User personas, feature priorities, competitive advantage

4. Book Writing Project

- Process: Exploration of outline/voice, then execution of chapter drafts
- Execution tasks: Draft chapter, revise, fact-check, polish
- Exploration: Target audience, chapter sequence, examples vs narrative

The adaptations:

Notice the same structure applies:

- **Exploration** always covers: audience, core positioning, key constraints
- **Execution** always covers: the actual work, quality standards, task sequencing
- **Progress** always covers: completion milestones, decisions made, momentum tracking

The variables are *what* you're building, but the *process* is identical.

How to adapt an example to your context:

1. **Replace the domain:** Sarah's "target audience = remote workers" becomes your "target audience = enterprise teams" or "academic researchers"
2. **Replace the decisions:** Sarah's "membership model = rejected" becomes your equivalent closed decision

3. **Replace the tasks:** Sarah's "define pricing model" becomes your "define feature prioritization" or "plan curriculum"
4. **Keep the structure:** Four files, session rhythm, archive pattern—all the same

When you read an example in this book, ask: "What's the principle here, and how does it apply to my project?" Usually, the principle transfers even if the details are completely different.

The Meta-Practice

Here's something worth knowing: This entire book is an example of the system working.

The book you're reading was written using the Context Kit. The author used:

- **Process:** To track which chapter was active
- **Exploration:** To track decisions about structure, tone, audience, examples
- **Execution:** To maintain guidelines for style and a checklist of chapters
- **Progress:** To log what was completed each session

Every chapter got done because the next chapter was always visible in Execution. Decisions didn't get revisited because they were closed in Exploration. Sessions didn't waste time on context because Process was explicit.

And the AI helped maintain the files. At the end of each session, the author asked the AI to update all four files based on what got accomplished. The AI generated the updates, the author reviewed and saved them. This book is proof the workflow works.

The system works for books, coffee shops, products, renovations, and anything else that requires context across multiple sessions. Not because it's magic, but because it does one thing well: it makes what matters explicit so you can focus on the work.

Final Thought

You don't need to be perfect with this system. You don't need to update files every single session (though it helps). You don't need to follow the exact structure (adapt it).

What you need is: a way to remember what you decided, a record of what you did, clarity on what's next, and a consistent session structure.

The Context Kit provides all of that. The AI helps you maintain it.

Use it. Refine it. Make it yours. And come back to projects months later knowing exactly where you left off.

That's the whole point. That's the entire system.

You've Read the Book, Now What?

You understand the system. Here's how to start using it.

Step 1: Copy the templates. Go to Appendix A. Copy all four blank templates. Save them in a folder for your project: `your-project/kit/01-process.md` (and so on). You now have the structure.

Step 2: Start your first project. Pick something you're actually working on—not a test project, something real. Fill in the basics: What phase are you in? What's your current task? What are 2-3 key guidelines? Add one Progress entry with today's date. You don't need to be perfect. Just start.

Step 3: Work for three sessions. Use the system for three sessions. Upload your files, work with the AI, update the files at the end. Three sessions is enough to see if the rhythm works for you. You'll start to notice: less re-explaining, clearer focus, easier resuming.

Step 4: Refine what doesn't work. After three sessions, look at your files. Are your guidelines too vague? Is your task list too detailed? Is Current State unclear? Adjust. The system should serve you, not constrain you. Make it yours.

Step 5: Keep going. That's it. The system works through repetition. The more you use it, the more natural it becomes. A month from now, you won't think about the files—you'll just work. The context will be there when you need it.

The Context Kit isn't magic. It's just organized information that makes AI collaboration sustainable. You've learned the system. Now use it.

Stay in context.

Appendices

Appendix A: Template Library

Copy these blank templates to start your project. Customize the section names and structure to match your work.

Template: 01-process.md

```
# PROCESS

## Workflow Stages
1. **Exploration** (Exploring ideas and possibilities)
2. **Execution** (Defining guidelines, planning structure, doing the work)
3. **Progress** (Tracking what is done)

## Rules for AI
- Focus only on the active phase in Current State
- Do not revisit or change closed decisions unless explicitly asked
- Follow the naming standards and guidelines defined in Execution
- Prioritize clear structure and reusable outcomes

---

# CURRENT STATE

**Phase:** [Your current workflow phase]
**Last Completed:** [Your most recent milestone]
**Next:** [Your active goal]
```

Template: 02-exploration.md

```
# EXPLORATION

## Open Topics

**Topic Name:**
- Questions to resolve
```

- Options being considered
- Tensions not yet decided

Closed Decisions

****[YYYY-MM-DD] Decision Name****

- What was decided
- Why this was the best choice
- Relevant context or alternatives considered

Template: 03-execution.md

EXECUTION

Guidelines

****Style & Voice:****

- [Your tone, voice, or approach]

****Rules & Constraints:****

- [Specific requirements]
- [Things to avoid]

****Quality Bar:****

- [Standards for completion]

Plan

****Goals:****

- [What you're building or achieving]

****Structure:****

- [Phases, outline, or organizational approach]

****Out of Scope:****

- [What's explicitly not included]

****Done Looks Like:****

- [Criteria for completion]

Action

****Right Now:****

[Current focus]

```
**Tasks:**

### Phase 1: [Phase Name]
- [ ] Task
- [ ] Task
  - [ ] Subtask

### Phase 2: [Phase Name]
- [ ] Task
- [ ] Task
```

Template: 04-progress.md

```
# PROGRESS LOG

## YYYY-MM-DD
- **Completed:** [What got done] or **Decided:** [What decisions were made]
or **In Progress:** [What's underway] or **Next:** [What comes next]
```

How to Use These

1. Copy each template into a new file
2. Save them in a folder for your project (e.g., **my-project/kit/01-process.md**)
3. Fill in the sections with information specific to your project
4. Use the two-message upload pattern when you start a new session (see Chapter 5)

That's it. You're ready to work.

Template: 00-prompts.md (Optional)

Copy this template for quick access to your standard session prompts.

```
### Fresh Start / Resume Session
\\
I'm working on a project using a file system to keep context between
sessions. I'll share files in two messages:

**This message:** Process (how we work + where we are) + Execution
(guidelines + plan + tasks)
**Next message:** Exploration (what's been decided) + Progress (what's been
done)

Here's what I need from you:
```

- Focus only on what's in Current State
- Don't reopen closed decisions unless I ask
- Follow the guidelines and quality standards in Execution
- Help me update files at the end so I can pick this up later

I'll paste the first two files now.

```

---

### End of Session Update

```

Based on this session, please update my four files to reflect what we accomplished.

Generate the updated markdown for all four files.

```

---

### Context Reset (If AI Starts Drifting)

```

You're drifting from the guidelines. I've re-uploaded the Kit files. Use these as your sole source of truth and restart based only on what's in the files.

```

---

### Project-Specific Notes

Add any custom prompts or reminders specific to your project:

- Example: "Always ask about budget constraints before suggesting solutions"
- Example: "Use British spelling throughout"
- Example: "Reference the brand guidelines in Execution for tone"

## How to Use This

1. Save this as **00-prompts.md** in your project folder
2. Customize the Project-Specific Notes section for your needs
3. Copy prompts from this file when starting or ending sessions
4. Update prompts as your project evolves

This is your reference card—quick access to prompts you use regularly.

## Appendix B: Quick Start Guide

Get started in 6 steps.

### Step 1: Create Your Files (2 minutes)

Grab the templates from Appendix A. Create four files:

- `01-process.md`
- `02-exploration.md`
- `03-execution.md`
- `04-progress.md`

Save them in a folder for your project: `my-project/kit/`

### Step 2: Fill in the Basics (2 minutes)

#### In `01-process.md`:

- Current State: What phase are you in? (Exploration, Execution, or Progress)
- Last Completed: What was your last milestone?
- Next: What are you working on now?

#### In `03-execution.md`:

- Guidelines: What are 2-3 key rules for how you'll work?
- Plan: What are you building? What's the outline or phases?
- Right Now: What's the current task or focus?

#### In `02-exploration.md`:

- List 3-5 open questions you're thinking about

#### In `04-progress.md`:

- Add one entry: Today's date, and what you're starting

**Optional:** Create a `00-prompts.md` file with your standard session prompts for quick copy-paste access. See Appendix A for the template.

### Step 3: Upload and Start (1 minute)

Open a new chat with your AI. Paste this message:

I'm working on a project using a file system to keep context between

```
sessions.
I'll share files in two messages:

This message: Process (how we work + where we are) + Execution
(guidelines + plan + tasks)
Next message: Exploration (what's been decided) + Progress (what's been
done)

Here's what I need from you:
- Focus only on what's in Current State
- Don't reopen closed decisions unless I ask
- Follow the guidelines and quality standards in Execution
- Help me update files at the end so I can pick this up later

Here are my first two files:
```

Then paste `01-process.md` and `03-execution.md`.

In your next message, paste `02-exploration.md` and `04-progress.md`.

## Step 4: Work

Start working on your project. The AI knows the context.

## Step 5: End Your Session (5 minutes)

Before you close the chat:

1. **Update 03-execution.md:** Check off completed tasks. Add new ones if they come up.
2. **Update 02-exploration.md:** Move any closed decisions here with dates.
3. **Update 04-progress.md:** Add today's entry: what you completed, what you decided, what's next.
4. **Update 01-process.md:** Update Current State with where you are now.

Save your files.

## Step 6: Next Session

Open a new chat. Upload your files using the same pattern as Step 3.

You'll pick up exactly where you left off.

## That's It

Six steps, fifteen minutes total, and you have a system that works for months of work.

No lost context. No repeating yourself. No guessing where you left off.

Just work.

## Your First Session

Here's what the first 10 minutes actually look like:

**You:** [Open new chat, paste boot prompt, upload 01-process.md and 03-execution.md]

"I'm working on a project using a file system to keep context between sessions. I'll share files in two messages..."

[Paste the two files]

**AI:** "Got it. I can see you're in the Execution phase, working on menu pricing for your coffee shop. Your guidelines say to keep examples under 3 paragraphs and maintain a conversational tone. Your last completed milestone was finalizing your target demographic (remote workers, not students). What would you like to tackle first?"

**You:** [Second message, upload 02-exploration.md and 04-progress.md]

"Here's my context from earlier work."

**AI:** "Thanks. I can see you've already decided against student discounts and membership models for year one. You're targeting a 35% margin minimum with competitive pricing between \$4.50-6.50. Your Progress log shows you spent two sessions on location research. Ready to finalize those pricing tiers?"

**You:** "Yes. Let's create three pricing tiers: basic drip coffee, specialty drinks, and premium single-origin options."

**AI:** "Based on your 35% margin requirement and \$4.50-6.50 range, here's what I'd suggest..."

[Work continues]

That's it. The AI knows where you are, what you've decided, and what standards to follow. No re-explaining. No "wait, didn't we already talk about this?" You're working immediately.



## Appendix C: Troubleshooting Guide

### Common Pitfalls (And How to Avoid Them)

Most problems with the Context Kit aren't system failures—they're habits. Here are the mistakes people make and how to avoid them.

- **Pitfall 1:**

**Not updating files at session end** You finish working, close the chat, forget to update your files. Next session you can't remember if you finished that section or not. Your Current State is outdated. You waste 10 minutes reconstructing what happened.

**How to avoid it:** Make file updates part of closing the session, not an afterthought. Before you close the chat, ask the AI to update all four files. Review them, save them, done. Build the habit.

- **Pitfall 2:**

**Making Current State too vague** Your Process file says "Next: Continue working on the book." That's not helpful. Which chapter? Which section? What specifically are you doing?

**How to avoid it:** Be specific. "Next: Draft Chapter 5, file numbering clarification section" tells you and the AI exactly what's active. Vague Current State means confused next sessions.

- **Pitfall 3:**

**Archiving too aggressively** You move closed decisions to an archive because the Exploration file is getting long. Next session, the AI suggests ideas you already rejected because it can't see those closed decisions.

**How to avoid it:** Keep closed decisions active if they're still relevant to current work. Only archive decisions that are truly old and unrelated to your current phase. The history you still reference isn't archive material.

- **Pitfall 4:**

**Skipping the two-message pattern** You upload all four files in one message to save time. The upload fails or truncates. Or you upload just one file thinking "that's probably enough." The AI doesn't have full context.

**How to avoid it:** Use the two-message pattern every time. It's 30 seconds. Shortcuts create confusion.

- **Pitfall 5:**

**Not closing decisions explicitly** You and the AI discuss something and reach a conclusion, but you don't move it to Closed Decisions in Exploration. Three sessions later, the AI brings it up again because it doesn't know the door is closed.

**How to avoid it:** When you decide something, write it down. Date it. Explain why. Move it to Closed Decisions. Make implicit knowledge explicit.

- **Pitfall 6:**

**Treating guidelines as suggestions** Your Execution file says "examples under 3 paragraphs" but you keep writing 5-paragraph examples. Your guidelines drift. The AI follows what you do, not what you wrote.

**How to avoid it:** If you're consistently breaking a guideline, either follow it or change it. Don't let guidelines become aspirational—they should reflect how you actually work.

- **Pitfall 7:**

**Continuing sessions past the breaking point** Your chat is getting long. The AI is starting to drift. You keep pushing because you're "almost done." The session degrades. Your file updates are malformed.

**How to avoid it:** When you notice the AI struggling, stop. Update your files immediately, start a fresh chat, and resume. Pushing through context rot wastes more time than resetting.

## "The AI Keeps Suggesting Ideas I Already Rejected"

**Problem:** You've decided against something, but the AI keeps bringing it up.

**Fix:** Check your Exploration file. The rejected idea might still be in Open Topics instead of Closed Decisions. Move it to Closed Decisions with the date and reasoning.

When you upload Exploration in your next chat, the AI will see the closed decision and stop suggesting it.

**Prevention:** Always close decisions explicitly in Exploration. Don't just decide them verbally—write them down so the AI sees them.

## "I Lost Track of What I Was Working On"

**Problem:** You took a break and can't remember where you left off.

**Fix:** Check your Progress file. The last entry tells you exactly what you completed and what's next.

Open a new chat, upload your files, and jump back in. Your Progress file is your memory.

**Prevention:** Update your Progress file at the end of every session. Just one entry with today's date and what you did.

## "The Chat Filled Up Mid-Work"

**Problem:** You were in the middle of something and the chat stopped accepting messages.

**Fix:** Don't panic. Your files have your context. Open a new chat, upload your files using the two-message pattern, and keep working.

Update your files first (so you don't lose the current session's progress), then resume.

**Prevention:** Before your chat gets completely full, start a new one. When you notice messages are getting long, proactively open a fresh chat and resume. It's not a failure—it's normal project maintenance.

## "My Files Are Getting Too Long"

**Problem:** Your files have hundreds of entries or tasks, and they're hard to navigate.

**Fix:** Archive old content. Move completed tasks from Execution to a separate file called `03-execution-archive.md`. Move old Progress entries to a separate file called `04-progress-archive.md`.

Keep only active content in your main files. Archives are for history; main files are for now.

Your next session, upload only the active files.

## "I Started Working on Two Different Things and Got Confused"

**Problem:** You were working on Phase A, then jumped to Phase B, and now you're not sure what's active.

**Fix:** Check your Process file's Current State section. That should tell you what's in focus. If Current State is unclear, update it.

Your Execution file should also show which phase you're in and what the next task is.

**Prevention:** At the end of every session, update Current State explicitly. Write what you were working on and what comes next. This 30-second update prevents confusion later.

## "The AI Doesn't Remember Something I'm Sure We Discussed"

**Problem:** You reference something from a previous conversation and the AI has no record of it.

**Fix:** The AI can only know what's in the files you uploaded. If it's not in Process, Exploration, Execution, or Progress, the AI doesn't have access to it.

Check: Did you save this information to your files last session? If not, add it now. If you did save it, check that you uploaded the file in this chat.

**Prevention:** If something is important to your project, write it down in the files. Conversations disappear. Files persist.

## "I'm Stuck and Don't Know What to Do Next"

**Problem:** Your project has stalled. You're not sure what comes next.

**Fix:** Open your Execution file. The "Right Now" or next task should be listed there. If it's vague (like "Continue working on the outline"), make it more specific.

Ask the AI: "I'm stuck. What's the smallest, clearest next step I could take right now?" Work on that one thing. Often momentum restarts with just one small action.

Then update your files so the next session has a clear direction.

## "I Want to Revisit a Closed Decision"

**Problem:** You decided something weeks ago, but now you think it might need to change.

**Fix:** Open your Exploration file and find the closed decision. Read the original reasoning. Is it still valid?

If you decide to reopen it, move it back to Open Topics with a new entry: "Reopening because [reason]. Original decision: [what was decided and why]."

This documents why you're reconsidering it. Your next session will understand the context.

**Prevention:** When you close decisions, write down the reasoning. "We decided no student discount because margins are tight" is better than just "no student discount." Better reasoning helps you evaluate whether the decision should change later.

## "My Files Are Out of Sync With What I Actually Did"

**Problem:** You worked on something but forgot to update the files, and now they don't match reality.

**Fix:** Update them now. Check off the tasks you completed. Update Current State to match what you're actually working on. Add a Progress entry for this catch-up session.

It's not perfect, but it's better than letting the files decay further.

**Prevention:** Update files at the end of every session, before you close the chat. It's the most important habit. Everything else flows from this.

## "Should I Archive My Files or Keep Them Forever?"

**Problem:** You finished a project months ago and aren't sure whether to keep the files.

**Fix:** Keep them. They're your project's complete history and context. If you ever need to revisit, reference, or do a Phase 2, the files have everything.

You don't need to upload them to new chats (that project is closed), but don't delete them.

## "Can I Use This for Multiple Projects?"

**Problem:** You're managing several projects and aren't sure how to keep context separate.

**Fix:** Yes, absolutely. Give each project its own folder with its own set of four files:

```
project-a/
 kit/
 01-process.md
 02-exploration.md
 03-execution.md
 04-progress.md

project-b/
 kit/
 01-process.md
 02-exploration.md
 03-execution.md
 04-progress.md
```

When you work on Project A, upload A's files. When you switch to Project B, upload B's files. The AI never mixes them up.

## Still Stuck?

If your issue isn't here, apply the core principle: **Make implicit knowledge explicit.**

If something important isn't written down in your files, write it down now. If the AI doesn't know something, it's probably because it's not in the files you uploaded. If you're confused about what to do next, make Current State more explicit.

The system works because information is explicit and organized. When something breaks, usually it's because information got implicit again. Make it explicit, and it fixes itself.

## Advanced: Long-Term Archive Management

When you work on a project for months, files accumulate history. This section covers how to organize archives so they're useful references, not digital clutter.

### Archive Structure Pattern

Here's a proven directory structure for long-term projects:

```
project-name/
├── ... project files and folders
├── kit/
│ ├── 01-process.md
│ ├── 02-exploration.md
│ ├── 03-execution.md
│ └── 04-progress.md
└── archive/
 ├── progress-archive/
 │ ├── 2025-01-progress.md
 │ └── 2024-12-progress.md
 ├── execution-archive/
 │ ├── 2025-01-completed.md
 │ └── 2024-12-completed.md
 ├── exploration-archive/
 │ └── archived-decisions.md
 └── reference/
 └── related-projects.md
```

### Naming Conventions

**Date archives by completion month**, not by when you archived them. A file named **2025-01-progress.md** contains entries from January 2025, regardless of when you moved it to the archive.

**Be descriptive with content ranges**. Instead of just **progress-archive.md**, use **2025-01-progress-entries-01-30.md** to show exactly what's inside.

**Group by file type**, not arbitrary date ranges. Keep all Progress archives together, all Execution archives together. This makes finding specific history easier.

## What to Archive (and When)

**Progress entries:** Archive entries older than your current work phase, typically 4-6 weeks for active projects. If you're in month 3, archive months 1-2. Keep the current month active.

**Execution tasks:** Archive completed tasks once you have 20+ completed items cluttering your active file. Move the entire "completed" section to a dated archive file.

**Exploration decisions:** Only archive decisions truly unrelated to current work. If a decision still informs what you're doing now, keep it active. Old doesn't mean irrelevant.

**Process file:** Almost never needs archiving. It stays current by nature.

## What NOT to Archive

Don't archive active, relevant decisions just because they're old. Age isn't the criterion—relevance is.

Example: You decided on your target audience in month 1. You're now in month 4 working on marketing. That audience decision is still actively relevant. Don't archive it just because it's three months old.

## Using the AI to Create Archives

When it's time to archive, ask the AI to help:

I need to archive my completed tasks from the last two months.

Please create a separate archive file with those completed tasks (with their dates) and generate a cleaned Execution file with only active tasks remaining.

The AI will:

- Extract the completed items
- Create the archive file with proper dating
- Generate a clean active file
- Preserve all the information

You review both files, save them in the appropriate folders, and continue working with the lean active file.

## Archive Metadata

Each archived file should include a brief header explaining what it contains:

```
Progress Archive: January 2025
Date Range: 2025-01-01 to 2025-01-31
Entries: 1-30
Phase: Execution - Chapters 4-6 drafting
```

**\*\*Archived:\*\*** 2025-02-01

This metadata helps you find specific history months or years later.

## When Archives Become References

Sometimes archived content becomes reference material for new projects. You might archive a completed project's files, then reference those decisions when starting a similar project.

Keep these in the `archive/reference/` folder with clear naming:  
`coffee-shop-project-lessons.md` or `pricing-model-research.md`.

## Common issues and how to fix them.

### Detecting and Fixing Context Rot

#### The Warning Signs

You are likely experiencing context rot if the AI starts:

- **Breaking Rules:** Violating the tone or naming standards in **03-execution.md**.
- **Hallucinating History:** Suggesting ideas already marked as "Closed" in **02-exploration.md**.
- **Garbling Output:** Producing broken Markdown or repeating "Next Step" items.
- **Ignoring Files:** Favoring its own internal "memory" over the specific files you uploaded.

#### The Recovery Protocol

If you detect rot, do not keep working. Follow these steps to secure your project and reset the AI's focus:

- **Step 1: Secure the Files.** Immediately ask the AI for a final update of the **Current State** and **Action** list. Review these for errors, manually correct any "hallucinations," and save them to your local files. This ensures your progress is "safe" outside of the degrading chat.
- **Step 2: The Anchoring Prompt.** Re-upload the corrected files and point the AI back to the source: *"You are drifting. I have uploaded the latest Kit files. Use these as your sole source of truth and restart the task based only on those guidelines."*
- **Step 3: The Review Loop.** Ask the AI to summarize the **Current State** before it proceeds to ensure it has successfully "re-read" your new uploads.



- **Step 4: The Hard Reset.** If errors persist, copy your secured Kit files into a brand-new chat session and perform a fresh **System Boot**.

### Prevention Practices

- **Stay Lean:** Keep files under 300 lines; archive old data regularly.
  - **Session Boundaries:** Start a new chat for every major milestone in your **Action** list.
  - **Disable Platform Memory:** Turn off features like Gemini's "Saved Info" to ensure the AI stays project-specific.
-